

# Introduction

My code for the first question contains the following files: threadScheduling.c (Q1.1), processScheduling.c (Q1.2), a Makefile, and clean.sh. Since I wished to have the option to initiate threadScheduling and processScheduling separately, I chose not to run their .out files in the Makefile itself. Further, to clean up the remnants of past compilations, I made clean.sh which simply executes `rm -f *.out *.o`. As a general note, I realise that Nice values are inversely proportional to their priority as compared to other SCHED\_OTHER processes/threads. I chose to take tests in ascending order of Nice (for SCHED\_OTHER) and Priority (for realtime) values merely because I felt it would make the output I expected more apparent.

## 1.1. Thread Scheduling

threadScheduling.c works by initiating three threads using `pthread_create()` in the main function of the program. Each function is configured more or less in the same way:

1. The scheduling policy for the thread is defined using `pthread_setschedparam()`.
  - in the case of Thread A (allotted SCHED\_OTHER), it's nice value is then set using `setpriority()`.
2. A message stating that the thread has been created is printed.
3. The start time is stored using `clock_gettime()` in a timespec structure called start.
4. The counting function is executed.
5. Once the thread has finished executing, `clock_gettime` stores the end time in timespec stop.
6. The duration is calculated and printed.
7. The function returns to main, where `pthread_join()` is making the parent process wait.

Outputs were taken by averaging three runs of each priority value.

The expected output was for the realtime processes (B & C) to always finish after A, which was indeed the general pattern. Between B and C, the thread which finished first was completely incidental. Further, my testing showed that changing Nice values didn't seem to make much of a difference to runtimes, which can be explained by the fact that there were no other user programs running, and so, in a scarcity of other programs competing for CPU time, our thread gets executed first. Similarly, changing the priority Value for B & C also yielded very little difference.

## 1.2. Process Scheduling

processScheduling.c has a single function (`compile3()`) which executes the fork and exec functions required for creating all three processes. The processes needed to be initiated at the same-ish time, and therefore I've used a sort of nested fork based system to create all three processes. Start time is taken before each (so it is available to the parent process). If the `pid==0` (i.e. the current process is the child), `sched_setscheduler` is used to set the scheduler, and `execvp` is used to run a bash file corresponding to the process. Else, if `pid>0`, the process is the parent, and creation of the next process takes place and so on. At the end, a while loop makes the parent process wait for all three child processes to elapse, and stores the end time. Finally, the durations are calculated and printed. The bash files simply executes three separate (but same) copies of the kernel install files.

The results consistently show the realtime processes taking lesser time than the SCHED\_OTHER process, and that too by a large margin. This is unsurprising, given that, as mentioned in my analysis of the results of 1.1, the priority for those (1-99) is always greater than SCHED\_OTHER processes (0, and nice values are not comparable to the aforementioned priority), which are handled by the CFS.

## 2. Kernel Memory Copy

References:

1. <https://brennan.io/2016/11/14/kernel-dev-ep3/>
2. <https://cool-dev.in/posts/How-to-Implement-a-System-Call-In-Linux/>

I've created my syscall by adding the code present in `kernel_2d_memcpy.c` into `kernel/sys.c` in the extracted kernel files. Further, I modified `arch/x86/entry/syscalls/syscall_64.tbl` to include an entry for my required syscall.

The functionality for the syscall is simple: a buffer array is created in the kernel space, to which `copy_from_user()` copies the source array into, and `copy_to_user()` copies the buffer into the destination array.

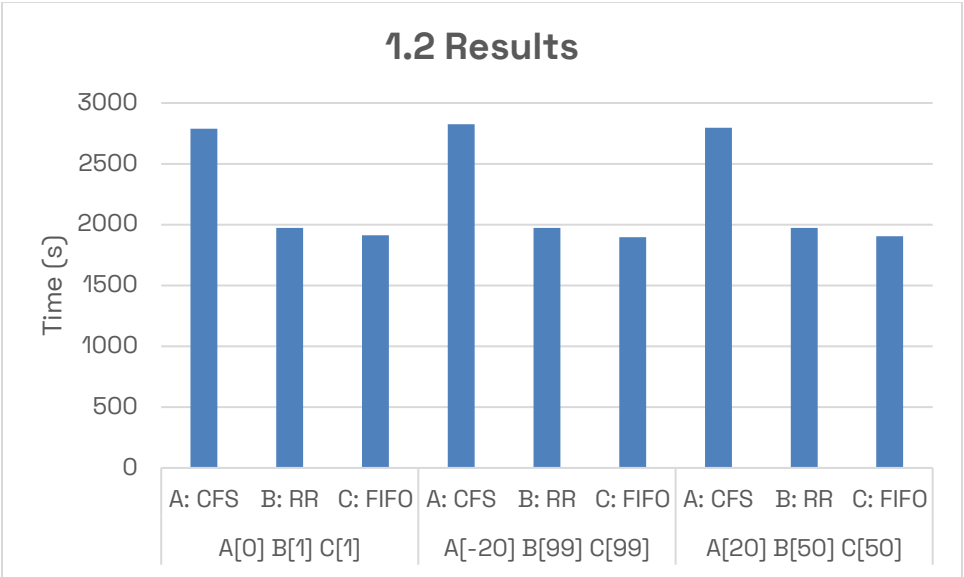
`testcall.c` is also quite simple – it simply creates an array and compares it manually to the copied array, which is generated using `syscall(syscall_number)`.

## Results



A[0] B[1] C[1]			A[-20] B[25] C[25]			A[-10] B[50] C[50]			A[10] B[75] C[75]			A[19] B[99] C[99]		
A: CFS	B: RR	C: FIFO	A: CFS	B: RR	C: FIFO	A: CFS	B: RR	C: FIFO	A: CFS	B: RR	C: FIFO	A: CFS	B: RR	C: FIFO
2.82	2.297	2.06166	3.8763	2.0673	2.1476	4.1366	2.0816	2.164	3.0983	2.260867	2.06733	3.14	2.31545	2.06962

## 1.2 Results



A[0] B[1] C[1]			A[-20] B[99] C[99]			A[20] B[50] C[50]		
A: CFS	B: RR	C: FIFO	A: CFS	B: RR	C: FIFO	A: CFS	B: RR	C: FIFO
2787.65	1974	1911.61	2824.147	1973.301	1897.775	2796.41	1973.72	1903.58